

LINEAR REGRESSION

DEBKANTA GHOSH

2026-02-04

1. Problem to demonstrate that the population regression line is fixed, but least square regression line varies.

```
rm(list=ls())
#STEP1: TAKE EQUIDISTANT POINT FROM 5 TO 10
x=seq(5,10,length.out=200)
x
## [1] 5.000000 5.025126 5.050251 5.075377 5.100503 5.125628 5.15075
## [8] 5.175879 5.201005 5.226131 5.251256 5.276382 5.301508 5.32663
## [15] 5.351759 5.376884 5.402010 5.427136 5.452261 5.477387 5.50251
## [22] 5.527638 5.552764 5.577889 5.603015 5.628141 5.653266 5.67839
## [29] 5.703518 5.728643 5.753769 5.778894 5.804020 5.829146 5.85427
## [36] 5.879397 5.904523 5.929648 5.954774 5.979899 6.005025 6.03015
## [43] 6.055276 6.080402 6.105528 6.130653 6.155779 6.180905 6.20603
## [50] 6.231156 6.256281 6.281407 6.306533 6.331658 6.356784 6.38191
## [57] 6.407035 6.432161 6.457286 6.482412 6.507538 6.532663 6.55778
## [64] 6.582915 6.608040 6.633166 6.658291 6.683417 6.708543 6.73366
## [71] 6.758794 6.783920 6.809045 6.834171 6.859296 6.884422 6.90954
## [78] 6.934673 6.959799 6.984925 7.010050 7.035176 7.060302 7.08542
## [85] 7.110553 7.135678 7.160804 7.185930 7.211055 7.236181 7.26130
## [92] 7.286432 7.311558 7.336683 7.361809 7.386935 7.412060 7.43718
## [99] 7.462312 7.487437 7.512563 7.537688 7.562814 7.587940 7.61306
## [106] 7.638191 7.663317 7.688442 7.713568 7.738693 7.763819 7.78894
## [113] 7.814070 7.839196 7.864322 7.889447 7.914573 7.939698 7.96482
```


```

## [120] 7.989950 8.015075 8.040201 8.065327 8.090452 8.115578 8.14070
4
## [127] 8.165829 8.190955 8.216080 8.241206 8.266332 8.291457 8.31658
3
## [134] 8.341709 8.366834 8.391960 8.417085 8.442211 8.467337 8.49246
2
## [141] 8.517588 8.542714 8.567839 8.592965 8.618090 8.643216 8.66834
2
## [148] 8.693467 8.718593 8.743719 8.768844 8.793970 8.819095 8.84422
1
## [155] 8.869347 8.894472 8.919598 8.944724 8.969849 8.994975 9.02010
1
## [162] 9.045226 9.070352 9.095477 9.120603 9.145729 9.170854 9.19598
0
## [169] 9.221106 9.246231 9.271357 9.296482 9.321608 9.346734 9.37185
9
## [176] 9.396985 9.422111 9.447236 9.472362 9.497487 9.522613 9.54773
9
## [183] 9.572864 9.597990 9.623116 9.648241 9.673367 9.698492 9.72361
8
## [190] 9.748744 9.773869 9.798995 9.824121 9.849246 9.874372 9.89949
7
## [197] 9.924623 9.949749 9.974874 10.000000

y=2+3*x
plot(x,y,type='l',lwd=3,main="plot population regresion function")

```

plot population regresion function



```

#step 2: Generate  $x_i(i = 1, 2, \dots, n)$  from  $Uniform(5, 10)$  and  $\epsilon_i(i = 1, 2, \dots, n)$  from  $N(0, 4)$ . Hence, compute  $y_1, y_2, \dots, y_n$ .

set.seed(123)
n=50
xi=runif(n,5,10)
ei=rnorm(n,0,4)
yi=2+3*xi+ei


#step3: Least square regression line
#two plot in same graph
plot(x,y,type='l',lwd=3,main="plot population regression line and sample regression line")
lin.reg=lm(yi~xi)
coef(lin.reg)

## (Intercept)           xi
## -0.09638929  3.30539569

abline(lin.reg,col="red",lwd=3)

```

Plot population regression line and sample regression line



```

#Step 4: Repeat steps 2-3 five times. Graph the 5 Least squares regression Lines over the population regression line obtained in Step 1.
set.seed(123)
x1=runif(n,5,10)
e1=rnorm(n,0,4)
y1=2+3*x1+e1

```

```

fit1=lm(y1~x1)
summary(fit1)

##
## Call:
## lm(formula = y1 ~ x1)
##
## Residuals:
##     Min      1Q  Median      3Q     Max 
## -9.0231 -2.2314 -0.2627  2.1970  8.7445 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) -0.09639   2.82610  -0.034   0.973    
## x1          3.30540   0.36519   9.051 5.96e-12 ***  
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
##
## Residual standard error: 3.761 on 48 degrees of freedom
## Multiple R-squared:  0.6306, Adjusted R-squared:  0.6229 
## F-statistic: 81.93 on 1 and 48 DF,  p-value: 5.962e-12 

x2=rnorm(n,5,10)
e2=rnorm(n,0,4)
y2=2+3*x2+e2
fit2=lm(y2~x2)

x3=rnorm(n,5,10)
e3=rnorm(n,0,4)
y3=2+3*x3+e3
fit3=lm(y3~x3)

x4=rnorm(n,5,10)
e4=rnorm(n,0,4)
y4=2+3*x4+e4
fit4=lm(y4~x4)

x5=rnorm(n,5,10)
e5=rnorm(n,0,4)
y5=2+3*x5+e5
fit5=lm(y5~x5)

#Data frame containing the five models' coefficients
coeff=data.frame(coef(fit1),
                  coef(fit2),
                  coef(fit3),
                  coef(fit4),
                  coef(fit5))

coeff

```

```


##           coef.fit1. coef.fit2. coef.fit3. coef.fit4. coef.fit5.
## (Intercept) -0.09638929   2.792188   1.392997   2.823089   2.032506
## x1          3.30539569   2.761042   3.073267   3.023608   3.028097

#Plot of the PRF and the Five SRFs

plot(x,y,type='l',lwd=3,main="Plot of the PRF and the Five SRFs")
lines(x1,predict(fit1),type='l',col="red")
lines(x2,predict(fit2),type='l',col="brown")
lines(x3,predict(fit3),type='l',col="green")
lines(x4,predict(fit4),type='l',col="blue")
lines(x5,predict(fit5),type='l',col="orange")
legend("topleft",legend=c("PRF","Model 1","Model 2","Model 3",
                         "Model 4","Model 5"),
       col=c("black","red","brown","green","blue","orange"),
       lwd=c(3,1,1,1,1,1))

```

Plot of the PRF and the Five SRFs



```

coeff

##           coef.fit1. coef.fit2. coef.fit3. coef.fit4. coef.fit5.
## (Intercept) -0.09638929   2.792188   1.392997   2.823089   2.032506
## x1          3.30539569   2.761042   3.073267   3.023608   3.028097

#Interpretation: PRF is fixed but SRF varies

```

Problem 2 Demonstrate that the estimators $\hat{\beta}_0$ and $\hat{\beta}_1$ minimize the Residual Sum of Squares (RSS).

```

rm(list=ls())

#step 1: Generate xi~U(5,10) of size 50, and ei~N(0,1)
#y=2+3x+e
n=50
set.seed(123)
x=runif(n,5,10)
xm=x-mean(x)
e=rnorm(n)
y=2+3*xm+e
fit=lm(y~xm)
summary(fit)

##
## Call:
## lm(formula = y ~ xm)
##
## Residuals:
##       Min     1Q Median     3Q    Max
## -2.25578 -0.55786 -0.06567  0.54926  2.18613
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept)  2.0562    0.1330   15.46   <2e-16 ***
## xm          3.0764    0.0913   33.70   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9404 on 48 degrees of freedom
## Multiple R-squared:  0.9594, Adjusted R-squared:  0.9586 
## F-statistic: 1135 on 1 and 48 DF,  p-value: < 2.2e-16

beta0=2.0562
beta=3.0764

#Creating the beta and beta0 grid values
beta0.grid=seq(-1,1,length.out=101)+beta0
beta.grid=seq(-1,1,length.out=101)+beta
#Computing RSS
RSS=c()
for(i in 1:length(beta.grid))
{
  RSS[i]=sum((y-beta0.grid[i]-beta.grid[i]*xm)^2)
}
RSS

## [1] 198.52541 192.34441 186.28828 180.35703 174.55065 168.86915 163.3125
2
## [8] 157.88076 152.57388 147.39187 142.33473 137.40247 132.59508 127.9125
6
## [15] 123.35492 118.92215 114.61425 110.43123 106.37308 102.43981 98.6314

```

```

0
## [22] 94.94788 91.38922 87.95544 84.64653 81.46249 78.40333 75.4690
5
## [29] 72.65963 69.97509 67.41542 64.98063 62.67071 60.48566 58.4254
9
## [36] 56.49019 54.67976 52.99421 51.43353 49.99772 48.68679 47.5007
3
## [43] 46.43954 45.50323 44.69179 44.00522 43.44353 43.00671 42.6947
7
## [50] 42.50769 42.44550 42.50817 42.69572 43.00814 43.44544 44.0076
1
## [57] 44.69465 45.50656 46.44335 47.50502 48.69155 50.00296 51.4392
5
## [64] 53.00040 54.68643 56.49734 58.43311 60.49376 62.67929 64.9896
8
## [71] 67.42495 69.98510 72.67012 75.48001 78.41477 81.47441 84.6589
2
## [78] 87.96831 91.40257 94.96170 98.64570 102.45458 106.38833 110.4469
6
## [85] 114.63046 118.93883 123.37208 127.93020 132.61319 137.42106 142.3538
0
## [92] 147.41141 152.59390 157.90126 163.33349 168.89060 174.57258 180.3794
3
## [99] 186.31116 192.36776 198.54924

df=data.frame(beta0.grid,beta.grid,RSS)
df

##      beta0.grid beta.grid      RSS
## 1      1.0562   2.0764 198.52541
## 2      1.0762   2.0964 192.34441
## 3      1.0962   2.1164 186.28828
## 4      1.1162   2.1364 180.35703
## 5      1.1362   2.1564 174.55065
## 6      1.1562   2.1764 168.86915
## 7      1.1762   2.1964 163.31252
## 8      1.1962   2.2164 157.88076
## 9      1.2162   2.2364 152.57388
## 10     1.2362   2.2564 147.39187
## 11     1.2562   2.2764 142.33473
## 12     1.2762   2.2964 137.40247
## 13     1.2962   2.3164 132.59508
## 14     1.3162   2.3364 127.91256
## 15     1.3362   2.3564 123.35492
## 16     1.3562   2.3764 118.92215
## 17     1.3762   2.3964 114.61425
## 18     1.3962   2.4164 110.43123
## 19     1.4162   2.4364 106.37308
## 20     1.4362   2.4564 102.43981
## 21     1.4562   2.4764  98.63140

```

## 22	1.4762	2.4964	94.94788
## 23	1.4962	2.5164	91.38922
## 24	1.5162	2.5364	87.95544
## 25	1.5362	2.5564	84.64653
## 26	1.5562	2.5764	81.46249
## 27	1.5762	2.5964	78.40333
## 28	1.5962	2.6164	75.46905
## 29	1.6162	2.6364	72.65963
## 30	1.6362	2.6564	69.97509
## 31	1.6562	2.6764	67.41542
## 32	1.6762	2.6964	64.98063
## 33	1.6962	2.7164	62.67071
## 34	1.7162	2.7364	60.48566
## 35	1.7362	2.7564	58.42549
## 36	1.7562	2.7764	56.49019
## 37	1.7762	2.7964	54.67976
## 38	1.7962	2.8164	52.99421
## 39	1.8162	2.8364	51.43353
## 40	1.8362	2.8564	49.99772
## 41	1.8562	2.8764	48.68679
## 42	1.8762	2.8964	47.50073
## 43	1.8962	2.9164	46.43954
## 44	1.9162	2.9364	45.50323
## 45	1.9362	2.9564	44.69179
## 46	1.9562	2.9764	44.00522
## 47	1.9762	2.9964	43.44353
## 48	1.9962	3.0164	43.00671
## 49	2.0162	3.0364	42.69477
## 50	2.0362	3.0564	42.50769
## 51	2.0562	3.0764	42.44550
## 52	2.0762	3.0964	42.50817
## 53	2.0962	3.1164	42.69572
## 54	2.1162	3.1364	43.00814
## 55	2.1362	3.1564	43.44544
## 56	2.1562	3.1764	44.00761
## 57	2.1762	3.1964	44.69465
## 58	2.1962	3.2164	45.50656
## 59	2.2162	3.2364	46.44335
## 60	2.2362	3.2564	47.50502
## 61	2.2562	3.2764	48.69155
## 62	2.2762	3.2964	50.00296
## 63	2.2962	3.3164	51.43925
## 64	2.3162	3.3364	53.00040
## 65	2.3362	3.3564	54.68643
## 66	2.3562	3.3764	56.49734
## 67	2.3762	3.3964	58.43311
## 68	2.3962	3.4164	60.49376
## 69	2.4162	3.4364	62.67929
## 70	2.4362	3.4564	64.98968
## 71	2.4562	3.4764	67.42495

```

## 72      2.4762    3.4964   69.98510
## 73      2.4962    3.5164   72.67012
## 74      2.5162    3.5364   75.48001
## 75      2.5362    3.5564   78.41477
## 76      2.5562    3.5764   81.47441
## 77      2.5762    3.5964   84.65892
## 78      2.5962    3.6164   87.96831
## 79      2.6162    3.6364   91.40257
## 80      2.6362    3.6564   94.96170
## 81      2.6562    3.6764   98.64570
## 82      2.6762    3.6964   102.45458
## 83      2.6962    3.7164   106.38833
## 84      2.7162    3.7364   110.44696
## 85      2.7362    3.7564   114.63046
## 86      2.7562    3.7764   118.93883
## 87      2.7762    3.7964   123.37208
## 88      2.7962    3.8164   127.93020
## 89      2.8162    3.8364   132.61319
## 90      2.8362    3.8564   137.42106
## 91      2.8562    3.8764   142.35380
## 92      2.8762    3.8964   147.41141
## 93      2.8962    3.9164   152.59390
## 94      2.9162    3.9364   157.90126
## 95      2.9362    3.9564   163.33349
## 96      2.9562    3.9764   168.89060
## 97      2.9762    3.9964   174.57258
## 98      2.9962    4.0164   180.37943
## 99      3.0162    4.0364   186.31116
## 100     3.0362    4.0564   192.36776
## 101     3.0562    4.0764   198.54924

df[which.min(RSS),]

##      beta0.grid beta.grid      RSS
## 51      2.0562    3.0764  42.4455

```

#We can verify the LSE minimises the RSS

3. Problem to demonstrate that least square estimators are unbiased

```

rm(list=ls())

#Step 1: Generate xi~U(5,10) of size 50, and ei~N(0,1)
n=50
set.seed(123)
x=runif(n,0,1)
e=rnorm(n)
y=2+3*x+e

```

```

#Step 2: obtain LSE
fit=lm(y~x)
summary(fit)

##
## Call:
## lm(formula = y ~ x)
##
## Residuals:
##       Min     1Q Median     3Q    Max 
## -2.25578 -0.55786 -0.06567  0.54926  2.18613 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept)  1.8576    0.2721   6.827 1.36e-08 ***
## x            3.3817    0.4565   7.408 1.75e-09 ***  
## ---    
## Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9404 on 48 degrees of freedom
## Multiple R-squared:  0.5334, Adjusted R-squared:  0.5237 
## F-statistic: 54.88 on 1 and 48 DF,  p-value: 1.745e-09

beta0.hat=1.8576
beta.hat=3.3817
#Repeat R=1000 times
set.seed(123)
R=10000
beta0=numeric(R)
beta=numeric(R)
n=50
for(i in 1:R) {
  x=runif(n,0,1)
  e=rnorm(n,0,1)
  y=2+3*x+e
  fit=lm(y~x)
  beta0[i]=coef(fit)[1]
  beta[i]=coef(fit)[2]
}
mean(beta0)

## [1] 2.005203

mean(beta)

## [1] 2.991338

#comment: the Long run means of beta0.hat and beta hat converges to true beta
#and beta0

var(beta0)

```

```

## [1] 0.08292927
var(beta)
## [1] 0.2495031
#Comment: var approaches 0 as R goes up. Hence LSE is consistent.

```

4.Comparing several simple linear regressions Attach “Boston” data from MASS library in R. Select median value of owneroccupied

```

rm(list=ls())
# Load required Library and data
library(MASS)
data(Boston)
# (a) Run separate simple Linear regressions
model_crime = lm(medv ~ crime, data = Boston)
model_nox = lm(medv ~ nox, data = Boston)
model_black=lm(medv ~ black, data = Boston)
model_lstat=lm(medv ~ lstat, data = Boston)

# Present output in a single table
summary_table=data.frame(
  Predictor = c("crime", "nox", "black", "lstat"),
  Coefficient = c(
    coef(model_crime)[2],
    coef(model_nox)[2],
    coef(model_black)[2],
    coef(model_lstat)[2]
  ),
  R_squared = c(
    summary(model_crime)$r.squared,
    summary(model_nox)$r.squared,
    summary(model_black)$r.squared,
    summary(model_lstat)$r.squared
  )
)
summary_table

##      Predictor  Coefficient R_squared
## crime      crime -0.41519028 0.1507805
## nox        nox -33.91605501 0.1826030
## black      black  0.03359306 0.1111961
## lstat      lstat -0.95004935 0.5441463

# (b) Identify the best model based on R-squared
best_model=summary_table[which.max(summary_table$R_squared), ]
best_model

```

```

## Predictor Coefficient R_squared
## lstat      lstat -0.9500494 0.5441463

# (c) Compare coefficients and usefulness of predictors
summary(model_crim)

##
## Call:
## lm(formula = medv ~ crim, data = Boston)
##
## Residuals:
##    Min     1Q   Median     3Q    Max
## -16.957 -5.449 -2.007  2.512 29.800
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 24.03311  0.40914  58.74 <2e-16 ***
## crim        -0.41519  0.04389  -9.46 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.484 on 504 degrees of freedom
## Multiple R-squared:  0.1508, Adjusted R-squared:  0.1491
## F-statistic: 89.49 on 1 and 504 DF,  p-value: < 2.2e-16

summary(model_nox)

##
## Call:
## lm(formula = medv ~ nox, data = Boston)
##
## Residuals:
##    Min     1Q   Median     3Q    Max
## -13.691 -5.121 -2.161  2.959 31.310
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 41.346     1.811   22.83 <2e-16 ***
## nox        -33.916     3.196  -10.61 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.323 on 504 degrees of freedom
## Multiple R-squared:  0.1826, Adjusted R-squared:  0.181
## F-statistic: 112.6 on 1 and 504 DF,  p-value: < 2.2e-16

summary(model_black)

##
## Call:
## lm(formula = medv ~ black, data = Boston)

```

```

## 
## Residuals:
##   Min     1Q Median     3Q    Max
## -18.884 -4.862 -1.684  2.932 27.763
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 10.551034  1.557463  6.775 3.49e-11 ***
## black       0.033593  0.004231  7.941 1.32e-14 ***
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 8.679 on 504 degrees of freedom
## Multiple R-squared:  0.1112, Adjusted R-squared:  0.1094 
## F-statistic: 63.05 on 1 and 504 DF,  p-value: 1.318e-14

summary(model_lstat)

## 
## Call:
## lm(formula = medv ~ lstat, data = Boston)
## 
## Residuals:
##   Min     1Q Median     3Q    Max
## -15.168 -3.990 -1.318  2.034 24.500
## 
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 34.55384  0.56263  61.41  <2e-16 ***
## lstat       -0.95005  0.03873 -24.53  <2e-16 ***
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
## Residual standard error: 6.216 on 504 degrees of freedom
## Multiple R-squared:  0.5441, Adjusted R-squared:  0.5432 
## F-statistic: 601.6 on 1 and 504 DF,  p-value: < 2.2e-16

```